

Bayesian regression for predicting combined cycle power plant output

Khang Nguyen, khang.nguyen@aalto.fi Son Le, son.le@aalto.fi

- 1 Overview
- 2 Linear model
- 3 Generalized additive model (GAM)
- 4 Model comparison
- 5 Conclusion
- 6 Appendix

- 1 Overview
- 2 Linear model
- 3 Generalized additive model (GAM)
- 4 Model comparison
- 5 Conclusion
- 6 Appendix

Introduction

- Thermodynamic systems are hard to construct and solve
- Combined cycle power plants (CCPPs) are examples
- Predict electrical output of CCPPs to maximize profit
- Past attempts: mathematical nonlinear equations, machine learning models
- Answer this problem with Bayesian approach

Data description





- Taken from the University of California Irvine Machine Learning Repository
- 9568 x 5 dataset, with AT (ambient temperature), V (exhaust vacuum), AP (atmospheric pressure), RH (relative humidity), and PE (full-load electrical output)

AT	V	AP	RH	PE
288.1	417.6	1024.1	731.7	463.3
298.3	629.6	1020.0	590.8	444.4
278.3	394.0	1012.2	921.4	488.6
294.0	573.2	1010.2	766.4	446.5
284.0	375.0	1009.2	966.2	473.9

Data description



Data description



Data analysis problem

- Fit regression models to predict the dependent variable by using independent variables
- Full-load electrical power output PE is predicted by the ambient temperature AT, the atmospheric pressure AP, the relative humidity RH, and the exhaust vacuum V
- Created two models to solve this problem: linear model and generalized additive model (GAM)

- 1 Overview
- 2 **Linear model**
- 3 Generalized additive model (GAM)
- 4 Model comparison
- 5 Conclusion
- 6 Appendix

Model description

- Implemented as the baseline for this analysis
- Has the form

$$y_i = \mathbf{x}_i^T \mathbf{w} + \epsilon_i, \quad i = 1, \dots, n$$
$$\epsilon_i \sim N(0, \sigma)$$

- Chosen priors

$$\mathbf{w} \sim N(0, 5)$$

$$\text{Intercept} \sim N(450, 50)$$

$$\sigma \sim \text{exponential}(0.05)$$

Convergence diagnostics


- Model built with `brms` and inference run with Stan
- No iteration exceeded maximum tree depth of 10
- Other HMC-specific convergence diagnostics were good
- \hat{R} -values were approximately 1
- n_{eff}/N ratios were good

Posterior predictive checks


- How well do the models describe the data?
- Simulate the data used for posterior predictive checks by sampling from the posterior predictive distribution:

$$p(\tilde{y} | y) = \int p(\tilde{y} | \theta) p(\theta | y) d\theta$$

Posterior predictive checks - kernel density estimates



Posterior predictive checks - error scatter plot



Sensitivity analysis

- 2 alternative set of priors were tested
- Uses only weakly informative priors, including the Intercept:

$$\text{Intercept} \sim \text{student}T(3, 0, 1)$$

- Different set of priors with domain knowledge for Intercept:

$$\text{Intercept} \sim N(450, 50)$$

- Conclusion: priors built from domain knowledge perform better.

- 1 Overview
- 2 Linear model
- 3 Generalized additive model (GAM)**
- 4 Model comparison
- 5 Conclusion
- 6 Appendix

Model description

- Captures nonlinear relationships through linear combination of smooth functions of explanatory variables
 - A smooth function has a linear basis expansion in some chosen basis
 - Splines are piecewise polynomials
 - Knots connect the pieces
 - Thin plate spline basis for smooth functions - avoids knot selection
- Formula:

$$\mu = \text{Intercept} + s(x_1) + s(x_2) + s(x_3) + s(x_4)$$

$$y \sim N(\mu, \sigma)$$

- Each smooth term $s(x_i)$ has 2 parts: fixed effect and random effect

Prior

- Standard deviation of each smooth term:

$$\sigma_i \sim \text{exponential}(0.1), \quad i = 1, \dots, 4$$

- Fixed effects:

$$\beta_i \sim N(0, 100), \quad i = 1, \dots, 4$$

- Standard deviation of errors:

$$\sigma \sim \text{exponential}(0.05)$$


- Intercept:

$$\text{Intercept} \sim N(460, 75)$$


Convergence diagnostics

- Model built with `brms` and inference run with Stan
- 3 iterations exceeded maximum tree depth of 10
- Other HMC-specific convergence diagnostics were good
- \hat{R} -values were approximately 1
- n_{eff}/N ratios were good

Posterior predictive checks - kernel density estimates



Posterior predictive checks - error scatter plot



Sensitivity analysis

- Two extra set of priors were tested
- Wider prior for standard deviations σ_i of smooth terms:

$$\sigma_i \sim \text{exponential}(0.001), \quad i = 1, \dots, 4$$

- Narrower priors for σ_i and fixed effects:

$$\sigma_i \sim \text{exponential}(1), \quad i = 1, \dots, 4$$


$$\beta_i \sim N(0, 1), \quad i = 1, \dots, 4$$

- Conclusion: priors for those parameters should not be too wide or too narrow.


- 1 Overview
- 2 Linear model
- 3 Generalized additive model (GAM)
- 4 Model comparison**
- 5 Conclusion
- 6 Appendix

Kernel density estimates

Linear model



GAM



Leave-one-out cross-validation comparison

- The difference in ELPD estimates of linear model and GAM:

	elpd_diff	se_diff	elpd_loo	se_elpd_loo	p_loo
gam	0.0	0.0	-21873.9	105.0	34.4
linear	-733.9	42.2	-22607.7	88.5	7.1

⇒ GAM is better.

Predictive performance - RMSE

- Linear model:

	Estimate	Est.Error	Q2.5	Q97.5
RMSEs	4.738	0.004	4.731	4.747

- GAM:

	Estimate	Est.Error	Q2.5	Q97.5
RMSEs	4.405	0.008	4.391	4.421

Predictive performance - Bayesian R2

- Linear model:

	Estimate	Est.Error	Q2.5	Q97.5
R2	0.91737	0.00051	0.91636	0.91833

- GAM:

	Estimate	Est.Error	Q2.5	Q97.5
R2	0.92911	0.00044	0.92823	0.92995

- 1 Overview
- 2 Linear model
- 3 Generalized additive model (GAM)
- 4 Model comparison
- 5 Conclusion**
- 6 Appendix

Conclusion

- Predict the full-load electrical output of a combined cycle power plant with conditions of the plant
- 2 Bayesian models with our prior belief
- GAM perform significantly better than linear model (ELPD, RMSE, Bayesian R2, and posterior predictive density checking)
- Possible improvements: variable selection and hierarchical model with groups created by an external clustering method

- 1 Overview
- 2 Linear model
- 3 Generalized additive model (GAM)
- 4 Model comparison
- 5 Conclusion
- 6 Appendix**

Variable conversion

```
dat_modified <-  
  dat %>% mutate(  
    V = V*10, # from cmHg to mmHg  
    RH = RH*10, # from percent to per-mille (‰)  
    AT = AT %>% conv_unit("C", "K") # from Celsius to Kelvin  
  )
```

Linear model - Stan code

```
// generated with brms 2.14.4
functions {
}
data {
  int<lower=1> N;  // total number of observations
  vector[N] Y;  // response variable
  int<lower=1> K;  // number of population-level effects
  matrix[N, K] X;  // population-level design matrix
  int prior_only;  // should the likelihood be ignored?
}
transformed data {
  int Kc = K - 1;
  matrix[N, Kc] Xc;  // centered version of X without an intercept
  vector[Kc] means_X;  // column means of X before centering
  for (i in 2:K) {
    means_X[i - 1] = mean(X[, i]);
    Xc[, i - 1] = X[, i] - means_X[i - 1];
  }
}
```


Linear model - Stan code

```
parameters {  
  vector[Kc] b; // population-level effects  
  real Intercept; // temporary intercept for centered predictors  
  real<lower=0> sigma; // residual SD  
}  
transformed parameters {  
}  
model {  
  // likelihood including all constants  
  if (!prior_only) {  
    target += normal_id_glm_lpdf(Y | Xc, Intercept, b, sigma);  
  }  
  // priors including all constants  
  target += normal_lpdf(b | 0, 5);  
  target += normal_lpdf(Intercept | 450, 50);  
  target += exponential_lpdf(sigma | 0.05)  
}
```


Linear model - Stan code

```
generated quantities {  
  // actual population-level intercept  
  real b_Intercept = Intercept - dot_product(means_X, b);  
  // additionally draw samples from priors  
  real prior_b = normal_rng(0,5);  
  real prior_Intercept = normal_rng(450,50);  
  real prior_sigma = exponential_rng(0.05);  
  // use rejection sampling for truncated priors  
  while (prior_sigma < 0) {  
    prior_sigma = exponential_rng(0.05);  
  }  
}
```


Linear model - Convergence diagnostics

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: PE ~ AT + V + AP + RH
## Data: train (Number of observations: 7724)
## Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##           total post-warmup samples = 4000
##
## Population-Level Effects:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept  1003.21    13.70   976.68  1029.96 1.00    1894    1969
## AT         -1.98     0.02    -2.02   -1.95 1.00    2355    2121
## V          -0.02     0.00    -0.02   -0.02 1.00    3209    2901
## AP          0.06     0.01     0.03    0.08 1.00    2189    1875
## RH         -0.02     0.00    -0.02   -0.01 1.00    3192    3220
##
## Family Specific Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma      4.52      0.04     4.45    4.59 1.00    2496    2292
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```


Linear model - Convergence diagnostics




Linear model - Convergence diagnostics



Linear model - Posterior predictive checks, error histograms



Linear model - Posterior predictive checks, LOO-PIT



Linear model - LOO-CV

```
##  
## Computed from 4000 by 7724 log-likelihood matrix  
##  
##           Estimate      SE  
## elpd_loo -22607.7   88.5  
## p_loo      7.1     0.7  
## looic      45215.5 176.9  
## -----  
## Monte Carlo SE of elpd_loo is 0.1.  
##  
## All Pareto k estimates are good ( $k < 0.5$ ).  
## See help('pareto-k-diagnostic') for details.
```


GAM - Stan code

```
// generated with brms 2.14.4
functions {
}
data {
  int<lower=1> N; // total number of observations
  vector[N] Y; // response variable
  // data for splines
  int Ks; // number of linear effects
  matrix[N, Ks] Xs; // design matrix for the linear effects
  // data for spline s(AT)
  int nb_1; // number of bases
  int knots_1[nb_1]; // number of knots
  // basis function matrices
  matrix[N, knots_1[1]] Zs_1_1;
  // data for spline s(V)
  int nb_2; // number of bases
  int knots_2[nb_2]; // number of knots
  // basis function matrices
  matrix[N, knots_2[1]] Zs_2_1;
```

GAM - Stan code

```
// data for spline s(AP)
int nb_3; // number of bases
int knots_3[nb_3]; // number of knots
// basis function matrices
matrix[N, knots_3[1]] Zs_3_1;
// data for spline s(RH)
int nb_4; // number of bases
int knots_4[nb_4]; // number of knots
// basis function matrices
matrix[N, knots_4[1]] Zs_4_1;
int prior_only; // should the likelihood be ignored?
}
transformed data {
}
```

GAM - Stan code

```

parameters {
  real Intercept; // temporary intercept for centered predictors
  vector[Ks] bs; // spline coefficients
  // parameters for spline s(AT)
  // standarized spline coefficients
  vector[knots_1[1]] zs_1_1;
  real<lower=0> sds_1_1; // standard deviations of spline coefficients
  // parameters for spline s(V)
  // standarized spline coefficients
  vector[knots_2[1]] zs_2_1;
  real<lower=0> sds_2_1; // standard deviations of spline coefficients
  // parameters for spline s(AP)
  // standarized spline coefficients
  vector[knots_3[1]] zs_3_1;
  real<lower=0> sds_3_1; // standard deviations of spline coefficients
  // parameters for spline s(RH)
  // standarized spline coefficients
  vector[knots_4[1]] zs_4_1;
  real<lower=0> sds_4_1; // standard deviations of spline coefficients
  real<lower=0> sigma; // residual SD
}

```

GAM - Stan code

```

transformed parameters {
  // actual spline coefficients
  vector[knots_1[1]] s_1_1;
  // actual spline coefficients
  vector[knots_2[1]] s_2_1;
  // actual spline coefficients
  vector[knots_3[1]] s_3_1;
  // actual spline coefficients
  vector[knots_4[1]] s_4_1;
  // compute actual spline coefficients
  s_1_1 = sds_1_1 * zs_1_1;
  // compute actual spline coefficients
  s_2_1 = sds_2_1 * zs_2_1;
  // compute actual spline coefficients
  s_3_1 = sds_3_1 * zs_3_1;
  // compute actual spline coefficients
  s_4_1 = sds_4_1 * zs_4_1;
}
generated quantities {
  // actual population-level intercept
  real b_Intercept = Intercept;
}

```

GAM - Stan code

```


model {
  // likelihood including all constants
  if (!prior_only) {
    // initialize linear predictor term
    vector[N] mu = Intercept + rep_vector(0.0, N) +
      Xs * bs + Zs_1_1 * s_1_1 + Zs_2_1 * s_2_1 +
      Zs_3_1 * s_3_1 + Zs_4_1 * s_4_1;
    target += normal_lpdf(Y | mu, sigma);
  }
  // priors including all constants
  target += normal_lpdf(Intercept | 460, 75);
  target += normal_lpdf(bs | 0, 100)
    - 4 * normal_lccdf(0 | 0, 100);
  target += exponential_lpdf(sds_1_1 | 0.1);
  target += std_normal_lpdf(zs_1_1);
  target += exponential_lpdf(sds_2_1 | 0.1);
  target += std_normal_lpdf(zs_2_1);
  target += exponential_lpdf(sds_3_1 | 0.1);
  target += std_normal_lpdf(zs_3_1);
  target += exponential_lpdf(sds_4_1 | 0.1);
  target += std_normal_lpdf(zs_4_1);
  target += exponential_lpdf(sigma | 0.05);
}

```


GAM - Convergence diagnostics

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: PE ~ s(AT) + s(V) + s(AP) + s(RH)
## Data: train (Number of observations: 7724)
## Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
## total post-warmup samples = 4000
##
## Smooth Terms:
##      Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sds(sAT_1)    14.36     4.14    8.38   24.45 1.00    1402    2044
## sds(sV_1)     19.79     5.21   11.93   32.27 1.00    1380    2046
## sds(sAP_1)     9.04     3.10    4.77   16.61 1.00    1386    2259
## sds(sRH_1)     6.19     3.08    1.92   13.51 1.00    1366    2126
##
## Population-Level Effects:
##      Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept    454.45      0.05  454.36  454.54 1.00     7461     2437
## sAT_1        -63.24    11.03  -83.87  -40.53 1.00     3462     2775
## sV_1          -4.68     9.96  -23.79   15.20 1.00     3319     2351
## sAP_1        -11.13     8.04  -26.79    4.56 1.00     2715     2393
## sRH_1         -1.22     8.26  -15.10   16.83 1.00     1682     2249
##
## Family Specific Parameters:
##      Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma      4.10      0.03    4.03    4.17 1.00     7378     2712
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```


GAM - Convergence diagnostics




GAM - Convergence diagnostics




GAM - Convergence diagnostics



GAM - Posterior predictive checks, error histograms



GAM - Posterior predictive checks, LOO-PIT



GAM - LOO-CV

```
##  
## Computed from 4000 by 7724 log-likelihood matrix  
##  
##           Estimate      SE  
## elpd_loo -21873.9 105.0  
## p_loo      34.4   2.1  
## looic      43747.8 210.1  
## -----  
## Monte Carlo SE of elpd_loo is 0.1.  
##  
## All Pareto k estimates are good (k < 0.5).  
## See help('pareto-k-diagnostic') for details.
```